

p0053r7 - C++ Synchronized Buffered Ostream

Lawrence Cowl, Peter Sommerlad, Nicolai Josuttis, Pablo Halpern

2017-11-07

Document Number:	p0053r7
Date:	2017-11-07
Project:	Programming Language C++
Audience:	LWG/LEWG

1 Introduction

At present, some stream output operations guarantee that they will not produce race conditions, but do not guarantee that the effect will be sensible. Some form of external synchronization is required. Unfortunately, without a standard mechanism for synchronizing, independently developed software will be unable to synchronize.

N3535 C++ Stream Mutexes proposed a standard mechanism for finding and sharing a mutex on streams. At the Spring 2013 standards meeting, the Concurrency Study Group requested a change away from a full mutex definition to a definition that also enabled buffering.

N3678 C++ Stream Guards proposed a standard mechanism for batching operations on a stream. That batching may be implemented as mutexes, as buffering, or some combination of both. It was the response to the Concurrency Study Group. A draft of that paper was reviewed in the Library Working Group, who found too many open issues on what was reasonably exposed to the 'buffering' part.

N3665 Uninterleaved String Output Streaming proposed making streaming of strings of length less than BUFSIZ appear uninterleaved in the output. It was a "minimal" functionality change to the existing standard to address the problem. The full Committee chose not to adopt that minimal solution.

N3750 C++ Ostream Buffers proposed an explicit buffering, at the direction of the general consensus in the July 2013 meeting of the Concurrency Study Group. In November 2014 a further updated version **N4187** was discussed in Library Evolution in Urbana and it was consensus to work with a library expert to get the naming and wording better suited to LWG expectations. Nico Josuttis volunteered to help the original authors. More information on the discussion is available at [LEWG wiki](#) and the corresponding [LEWG bug tracker entry \(20\)](#). This paper address issues raised with **N4187**. This paper has a change of title to reflect a change in class name, but contains the same fundamental approach.

1.1 Solution

We propose a `basic_osyncstream`, that buffers output operations for a wrapped stream. The `basic_osyncstream` will atomically transfer the contents of an internal stream buffer to a `basic_ostream`'s stream buffer on destruction of the `basic_osyncstream`.

The transfer on destruction simplifies the code and ensures at least some output in the presence of an exception.

The intent is that the `basic_osyncstream` is an automatic-duration variable with a relatively small scope which constructs the text to appear uninterleaved. For example,

```
{
    osyncstream bout(cout);
    bout << "Hello, " << "World!" << '\n';
}
```

or with a single output statement:

```
osyncstream{cout} << "The answer is " << 6*7 << endl;
```

2 Design

Here we list a set of previous objections to our proposal in the form of questions. We then give our reasons why other potential solutions do not work as well as our proposed solution.

Why can I not just use `cout`? It should be thread-safe.

You will not get data races for `cout`, but that is not true for most other streams. In addition there is no guarantee that output from different threads appears in any sensible order. Coherent output order is the main reason for this proposal.

Why must `osyncstream` be an `ostream`? Could a simple proxy wrapper work?

To support all existing and user-defined output operators, `osyncstream` must be an `ostream` subclass.

Can you make a flush of the `osyncstream` mean transfer the characters and flush the underlying stream?

No, because the point of the `osyncstream` is to collect text into an atomic unit. Flushes are often emitted in calls where the body is not visible, and hence unintentionally break up the text. Furthermore, there may be more than one flush within the lifetime of an `osyncstream`, which would impose a performance loss when an atomic unit of text needs only one flush. The design decision is only to flush the underlying stream if the `osyncstream` was flushed, and only once per atomic transfer of the character sequence. [p0053r5 introduced manipulators to allow both ways, but the default remains not to flush immediately.](#)

Can flush just transfer the characters and not flush the underlying stream?

The flush intends an effect on visible to the user. So, we should preserve at least one flush. The flush may not be visible to the code around the `osyncstream`, and so its programmer cannot do a manual flush, because attempting to flush the underlying stream that is shared among threads will introduce a data race.

Why do you specify `basic_syncbuf`? LWG and LEWG thought you wouldn't need it.

Users use the `streambuf` interface. Without access to the `basic_syncbuf` they would not be able call `emit()` on the underlying `streambuf` responsible the synchronization. Every stream in the standard is defined as a pair of `streambuf/stream`, for good reason. Ask Pablo Halpern if you need more to be convinced.

Where will the required lock/mutex be put? Will it be in every `streambuf` object changing the ABI?

That is one of the reasons why this must be put into the standard library. It is possible to implement with a global map from `streambuf*` to mutexes as the example code does, however, existing standard library implementations might have already a space for the mutexes (not breaking their ABI), because `cout/cerr` seem to require one and those might be the most important ones to wrap.

The design follows from the principles of the `iostream` library. If discussed a person knowledgeable about `iostream`'s implementation is favorable, because of its many legacy design decisions, that would no longer be taken by modern C++ class designers.

As with all existing stream classes, using a stream object or a `streambuf` object from multiple threads can result in a data race, this is also true for `osyncstream`. Its use enables sharing the wrapped `streambuf` object/output stream across several threads, if all concurrently using threads apply a local `osyncstream` around it.

We follow typical stream conventions of `basic_` prefixes and typedefs.

The constructor for `osyncstream` takes a non-const reference to a `basic_ostream` obtaining its stream buffer or a `basic_streambuf*` or can be put around such a stream buffer directly. Calling the `emit()` member function or destroying the `osyncstream` may write to the stream buffer obtained at construction.

The wording below permits implementation of `basic_osyncstream` with either a `stream_mutex` from N3535 or with implementations suitable for N3665, e.g. with Posix file locks [PSL]

2.1 Items to be discussed by LWG

- ~~Naming of the manipulators~~ (moved to separate paper p0753r0)
- What should be the delivery vehicle for this feature? I suggest both C++20 and the concurrency TS? **should be moved to C++-next working paper in Albuquerque.**

3 Wording

This wording is relative to the current C++ working draft. It could be integrated into a Concurrency TS accordingly.

Changes will happen in section 30 mainly.

In section 20.5.1.1 Library contents [contents] add an entry to table 16 (cpp.library.headers) for the new header `<syncstream>`.

3.1 30.1 General [input.output.general]

Insert a new entry in table 106 Input/output library summary (tab:iostreams.lib.summary)

| **30.x** Synchronized output stream <syncstream> |

3.2 30.3.1 Header <iosfwd> synopsis [iosfwd.syn]

Insert the following declarations to the synopsis of <iosfwd> in the namespace std.

```
template <class charT,
          class traits = char_traits<charT>,
          class Allocator = allocator<charT>>
    class basic_syncbuf;
using syncbuf = basic_syncbuf<char>;
using wsyncbuf = basic_syncbuf<wchar_t>;

template <class charT,
          class traits = char_traits<charT>,
          class Allocator = allocator<charT> >
    class basic_osyncstream;
using ostream = basic_osyncstream<char>;
using wostream = basic_osyncstream<wchar_t>;
```

3.3 30.x Synchronized output stream [syncstream]

Insert this new section 30.x in chapter 30 [input.output]

3.3.1 30.x.1 Header <syncstream> synopsis [syncstream.syn]

```
namespace std {
template <class charT,
          class traits,
          class Allocator>
    class basic_syncbuf;
using syncbuf = basic_syncbuf<char>;
using wsyncbuf = basic_syncbuf<wchar_t>;

template <class charT,
          class traits,
          class Allocator>
    class basic_osyncstream;
using ostream = basic_osyncstream<char>;
using wostream = basic_osyncstream<wchar_t>;
}
```

- ¹ The header <syncstream> provides a mechanism to synchronize execution agents writing to the same stream. It defines class templates `basic_osyncstream` and `basic_syncbuf`. The latter buffers output and transfer the buffered content into an object of type `basic_streambuf<charT, traits>` atomically with respect to such transfers by other `basic_syncbuf<charT, traits, Allocator>` objects referring to the same `basic_streambuf<charT, traits>` object. The transfer occurs when `emit()` is called and when the `basic_syncbuf<charT, traits, Allocator>` object is destroyed.

3.3.2 30.x.2 Class template basic_syncbuf [syncstream.syncbuf]

```

template <class charT,
          class traits,
          class Allocator>
class basic_syncbuf
    : public basic_streambuf<charT, traits> {

public:
    using char_type      = charT;
    using int_type       = typename traits::int_type;
    using pos_type       = typename traits::pos_type;
    using off_type       = typename traits::off_type;
    using traits_type     = traits;
    using allocator_type = Allocator;

    using streambuf_type = basic_streambuf<charT, traits>;

    explicit
    basic_syncbuf(streambuf_type* obuf = nullptr)
        : basic_syncbuf(obuf, Allocator()) { }
    basic_syncbuf(streambuf_type*, const Allocator&);
    basic_syncbuf(basic_syncbuf&&);
    ~basic_syncbuf();

    basic_syncbuf& operator=(basic_syncbuf&&);
    void swap(basic_syncbuf&);

    bool emit();
    streambuf_type* get_wrapped() const noexcept;
    allocator_type get_allocator() const noexcept;
    void          set_emit_on_sync(bool) noexcept;

protected:
    int sync() override;

private:
    streambuf_type* wrapped; // exposition only
    bool          emit_on_sync{}; // exposition only
};

template <class charT, class traits, class Allocator>
void swap(basic_syncbuf<charT, traits, Allocator>&,
          basic_syncbuf<charT, traits, Allocator>&);

```

3.4 30.x.2.1 basic_syncbuf constructors [syncstream.syncbuf.cons]

```
basic_syncbuf(streambuf_type* obuf, const Allocator& allocator);
```

- 1 *Effects:* Constructs the `basic_syncbuf` object and sets wrapped to `obuf` which will be the final destination of associated output.
- 2 *Remarks:* A copy of `allocator` is used to allocate memory for internal buffers holding the associated output.
- 3 *Throws:* Nothing unless constructing a mutex or allocating memory throws.
- 4 *Postconditions:* `get_wrapped() == obuf` && `get_allocator() == allocator`.

```
basic_syncbuf(basic_syncbuf&& other);
```

- 5 *Effects:* Move constructs from `other` (Table 23 [moveconstructable]).
- 6 *Postconditions:* The value returned by `this->get_wrapped()` is the value returned by `other.get_wrapped()` prior to calling this constructor. Output stored in `other` prior to calling this constructor will be stored in `*this` afterwards. `other.rdbuf()->pbase() == other.rdbuf()->pptr()` and `other.get_wrapped() == nullptr`.
- 7 *Remarks:* This constructor disassociates `other` from its wrapped stream buffer, ensuring destruction of `other` produces no output.

3.5 30.x.2.2 basic_syncbuf destructor [syncstream.syncbuf.dtor]

```
~basic_syncbuf();
```

- 1 *Effects:* Calls `emit()`.
- 2 *Throws:* Nothing. If an exception is thrown from `emit()`, that exception is caught and ignored.

3.5.1 30.x.2.3 basic_syncbuf assign and swap [syncstream.syncbuf.assign]

```
basic_syncbuf& operator=(basic_syncbuf&& rhs) noexcept;
```

- 1 *Effects:* Calls `emit()` then move assigns from `rhs`. After the move assignment `*this` has the observable state it would have had if it had been move constructed from `rhs` (see [syncstream.syncbuf.ctor]).
- 2 *Returns:* `*this`.
- 3 *Postconditions:* `rhs.get_wrapped() == nullptr`. If `allocator_traits<Allocator>::propagate_on_container_move_assignment::value == true`, then `this->get_allocator() == rhs.get_allocator()`; otherwise the allocator is unchanged.
- 4 *Remarks:* This assignment operator disassociates `rhs` from its wrapped stream buffer ensuring destruction of `rhs` produces no output.

```
void swap(basic_syncbuf& other) noexcept;
```

- 5 *Requires:* `allocator_traits<Allocator>::propagate_on_container_swap::value || this->get_allocator() == other.get_allocator()`.
- 6 *Effects:* Exchanges the state of `*this` and `other`.

```
template <class charT, class traits, class Allocator>
void swap(basic_syncbuf<charT, traits, Allocator>& a,
         basic_syncbuf<charT, traits, Allocator>& b) noexcept;
```

7 *Effects:* Equivalent to `a.swap(b)`.

3.5.2 30.x.2.4 basic_syncbuf member functions [syncstream.syncbuf.mfun]

```
bool emit();
```

1 *Effects:* Atomically transfers the contents of the internal buffer to the stream buffer `*wrapped`, so that they appear in the output stream as a contiguous sequence of characters. If and only if a call was made to `sync()` since the last call of `emit()`, `wrapped->pubsync()` is called.

2 *Returns:* `true` if all of the following conditions hold; otherwise `false`:

(2.1) — `wrapped != nullptr`.

(2.2) — All of the characters in the associated output were successfully transferred.

(2.3) — The call to `wrapped->pubsync()` (if any) succeeded.

3 *Postconditions:* On success the associated output is empty.

4 *Synchronization:* All `emit()` calls transferring characters to the same stream buffer object appear to execute in a total order consistent with *happens-before* where each `emit()` call *synchronizes-with* subsequent `emit()` calls in that total order.

5 *Remarks:* May call member functions of `wrapped` while holding a lock uniquely associated with `wrapped`.

```
streambuf_type* get_wrapped() const noexcept;
```

6 *Returns:* `wrapped`.

```
allocator_type get_allocator() const noexcept;
```

7 *Returns:* A copy of the allocator set in the constructor or from assignment.

```
void set_emit_on_sync(bool b) noexcept;
```

8 *Effects:* `emit_on_sync=b`.

3.5.3 30.x.2.5 basic_syncbuf overridden virtual member functions [syncstream.syncbuf.virtuals]

```
int sync() override;
```

1 *Effects:* Record that the wrapped stream buffer is to be flushed. Then, if `emit_on_sync == true`, calls `emit()`. [*Note:* If `emit_on_sync == false`, the actual flush is delayed until a call to `emit()`. — *end note*]

2 *Returns:* If `emit()` was called and returned `false`, returns `-1`; otherwise `0`.

3.6 30.x.3 Class template basic_osyncstream [syncstream.osyncstream]

```

template <class charT,
          class traits,
          class Allocator>
class basic_osyncstream
: public basic_ostream<charT, traits>
{
public:
    using char_type      = charT;
    using int_type       = typename traits::int_type;
    using pos_type       = typename traits::pos_type;
    using off_type       = typename traits::off_type;
    using traits_type     = traits;
    using allocator_type = Allocator;
    using streambuf_type = basic_streambuf<charT, traits>;
    using syncbuf_type   = basic_syncbuf<charT, traits, Allocator>;

    basic_osyncstream(streambuf_type*, const Allocator&);
    explicit basic_osyncstream(streambuf_type* obuf)
        : basic_osyncstream(obuf, Allocator()) { }
    basic_osyncstream(basic_ostream<charT, traits>& os, const Allocator& allocator)
        : basic_osyncstream(os.rdbuf(), allocator) { }
    explicit basic_osyncstream(basic_ostream<charT, traits>& os)
        : basic_osyncstream(os, Allocator()) { }
    basic_osyncstream(basic_osyncstream&&) noexcept;
    ~basic_osyncstream();
    basic_osyncstream& operator=(basic_osyncstream&&) noexcept;
    void emit();
    streambuf_type* get_wrapped() const noexcept;
    syncbuf_type* rdbuf() const noexcept { return &sb ; }
private:
    syncbuf_type sb; // exposition only
};

```

¹ `Allocator` shall meet the allocator requirements [allocator.requirements].

² [*Example:* Use a named variable within a block statement for streaming.

```

{
    osyncstream bout(cout);
    bout << "Hello, ";
    bout << "World!";
    bout << endl; // flush is noted
    bout << "and more!\n";
} // characters are transferred and cout is flushed

```

— *end example*]

³ [*Example:* Use a temporary object for streaming within a single statement. `cout` is not flushed.

```

osyncstream(cout) << "Hello, " << "World!" << '\n';

```


— end example]

3.6.1 30.x.3.1 basic_osyncstream constructors [syncstream.osyncstream.cons]

```
basic_osyncstream(streambuf_type* buf, const Allocator& allocator);
```

1 *Effects:* Initializes `sb` from `buf` and `allocator` and initializes the base class with `basic_ostream(&sb)`.

2 [*Note:* If wrapped stream buffer pointer refers to a user provided stream buffer then its implementation must be aware that its member functions might be called from `emit()` while a lock is held. — end note]

3 *Postconditions:* `get_wrapped() == buf`.

```
basic_osyncstream(basic_osyncstream&& other) noexcept;
```

4 *Effects:* Move constructs from `other`. This is accomplished by move constructing the base class, and the contained `basic_syncbuf sb`. Next `basic_ostream<charT, traits>::set_rdbuf(&sb)` is called to install the `basic_syncbuf sb`.

5 *Postconditions:* The value returned by `get_wrapped()` is the value returned by `os.get_wrapped()` prior to calling this constructor. `nullptr == other.get_wrapped()`.

3.6.2 30.x.3.2 basic_osyncstream destructor [syncstream.osyncstream.dtor]

```
~basic_osyncstream();
```

1 *Effects:* Calls `emit()`. If an exception is thrown from `emit()`, that exception is caught and ignored.

3.6.3 30.x.3.3 Assignment [syncstream.osyncstream.assign]

```
basic_osyncstream& operator=(basic_osyncstream&& rhs) noexcept;
```

1 *Effects:* First, calls `emit()`. If an exception is thrown from `emit()`, that exception is caught and ignored. Move assigns `sb` from `rhs.sb`. [*Note:* This disassociates `rhs` from its wrapped stream buffer ensuring destruction of `rhs` produces no output. — end note]

2 *Postconditions:* Primarily, `nullptr == rhs.get_wrapped()`. Also, `get_wrapped()` returns the value previously returned by `rhs.get_wrapped()`.

3.6.4 30.x.3.4 Member functions [syncstream.osyncstream.members]

```
void emit();
```

1 *Effects:* Calls `sb.emit()`. If this call returns `false`, calls `setstate(ios::badbit)`.

[*Example:* A flush on a `basic_osyncstream` does not flush immediately:

```
{
    osyncstream bout(cout);
    bout << "Hello," << '\n'; // no flush
    bout.emit(); // characters transferred; cout not flushed
    bout << "World!" << endl; // flush noted; bout not flushed
    bout.emit(); // characters transferred; cout flushed
```

```
    bout << "Greetings." << '\n'; // no flush
} // characters transferred; cout not flushed
```

— end example]

[*Example:* The function `emit()` can be used to catch exceptions from operations on the underlying stream.

```
{
    ostream bout(cout);
    bout << "Hello, " << "World!" << '\n';
    try {
        bout.emit();
    } catch ( ... ) {
        // stuff
    }
}
```

— end example]

```
streambuf_type* get_wrapped() const noexcept;
```

2 *Returns:* `sb.get_wrapped()`.

[*Example:* Obtaining the wrapped stream buffer with `get_wrapped()` allows wrapping it again with an `ostream`. For example,

```
{
    ostream bout1(cout);
    bout1 << "Hello, ";
    {
        ostream(bout1.get_wrapped()) << "Goodbye, " << "Planet!" << '\n';
    }
    bout1 << "World!" << '\n';
}
```

produces the *uninterleaved* output

```
Goodbye, Planet!
Hello, World!
```

— end example]

3.7 Feature test macro

For the purposes of SG10, we recommend the feature-testing macro name `__cpp_lib_syncstream`.

3.8 Implementation

An example implementation is available on https://github.com/PeterSommerlad/SC22WG21_Papers/tree/master/workspace/p0053_basic_ostreambuf

4 Revisions

Each section lists the revisions in the following version from the version given in the heading.

4.1 P0053R7

Final finishing touches to wording and layout based on LWG feedback. Added feature test macro

4.2 P0053R6

Revisions based on feedback from LWG in Toronto. Extracted the manipulators to p0753, so that they can be reconsidered by LEWG.

4.3 P0053R5

This paper was discussed by LWG in Toronto. All recommended changes have been incorporated into the next revision.

4.4 D0053R4 - P0053R4

This version was only published on the Kona Wiki. The manipulators were extracted into a separate paper to ease forward progress with this paper, even though the wording of the manipulator specification was already reviewed by LWG.

- Translate text to LaTeX
- Added manipulator support so that logging frameworks that get an `osyncstream` passed can rely on flushes happening.
- Make sure that temporary stream objects can be used safely (it is already in the standard, Pablo!)
- Ensured section/table numbers match current working draft (as of 06/2017 before the mailing)

4.5 P0053R3

- Takes input from Pablo Halpern and re-instantiate the stream buffer that performs the synchronization.
- Split the constructors with a defaulted allocator parameter to one single-argument one being explicit and one non-explicit taking 2 arguments.

4.6 P0053R2

- Remove the "may construct a mutex" notes.
- Remove "may destroy mutex" notes.
- Clarify `osyncstream` flush behavior in an example.
- Make minor editorial fixes.

4.7 P0053R1

- Provide a `typedef` for the wrapped stream buffer and use it to shorten the specification as suggested by Daniel Krügler.
- Provide move construction and move assignment and specify the moved-from state to be detached from the wrapped stream buffer.
- Rename `get()` to `rdbuf_wrapped()` and provide `noexcept` specification.

- Changed to explicitly rely on wrapping a stream buffer, instead of an `ostream` object and adjust explanations accordingly.

4.8 P0053R0

- Add remark to note that exchanging the stream buffer while the stream is wrapped causes undefined behavior and added a note to warn stream buffer implementers about the lock being held in `emit()`. Call `setstate(badbit)` if IO errors occur in `emit()`.
- Replace code references to `basic_streambuf` by the term stream buffer introduced in `[stream.buffer]`.
- Provide an example implementation.
- The lock is to be associate to the underlying `basic_streambuf` instead of the `basic_stream`.
- Added an `Allocator` constructor parameter.
- Moves destructor example to `emit()`.
- Clarifies wording about synchronization and flushing (several times).
- List the new header in corresponding table.
- Provide type aliases in `<iosfwd>`.
- Removed copy constructor in favor of providing `get()`.
- Notify that move construction and assignment is deleted.
- Moved class `noteflush_streambuf` into an implementation note.
- Add a design subsection that states that a header test is a sufficient feature test.

4.9 N4187

- Updated introduction with recent history.
- Rename `ostream_buffer` to `osyncstream` to reflect its appearance is more like a stream than like a buffer.
- Add an example of using `osyncstream` as a temporary object.
- Add an example of a `osyncstream` constructed with another `osyncstream`.
- Clarify the behavior of nested `osyncstream` executions.
- Clarify the behavior of exceptions occurring with the `osyncstream` destructor.
- Clarify the deferral of flush from the `osyncstream`'s `streambuf` to the final `basic_ostream`.
- Limit the number of references to `noteflush_stringbuf` in anticipation of the committee removing it from the specification.
- Rename `noteflush_stringbuf` to `noteflush_streambuf` to hide possible implementation details.
- Change the base class of `noteflush_streambuf` from `basic_stringbuf` to `basic_streambuf`.

4.10 N4069

- Added note to sync as suggested by BSI via email.

4.11 N3978

- Added a Design section.
- Clarify the reference capturing behavior of the `ostream_buffer` constructors.
- Added `noexcept` and `const` as appropriate to members.
- Added note on throwing wrapped streams.
- Change the `noteflush_stringbuf` public member variable `needsflush` to a public member query function `flushed`.
- Removed the public member function `noteflush_stringbuf::clear`.
- Minor synopsis formatting changes.
- Incorporated feedback from SG1 and Dietmar Kühl in specific in Rapperswil.

4.12 N3892

- Flush the `ostream` if and only if the `ostream_buffer` was flushed.
- Add the `clear_for_reuse` function.
- Change the design from inheriting from `basic_ostream` to using a `noteflush_stringbuf`, which is a slightly modified `basic_stringbuf`. The modification is to note the flush rather than act upon it.

4.13 N3750

- Change name to `basic_ostream_buffer` and add the usual typedefs.
- Change interface to inherit from `basic_ostringstream` rather than provide access to a member of that type.
- Add a Revisions section.